

Dimension Reduction Analysis: pbmc

Generated report #1

Seed: 0 | Plan revision: 1

Dataset profile & preprocessing

Peripheral blood mononuclear cells (PBMCs) from a healthy donor, profiled with 10x Genomics single-cell RNA sequencing (scRNA-seq). This is the standard "pbmc3k" dataset used throughout the Scanpy tutorials and widely used as a benchmark for single-cell analysis methods. Loaded via ``scanpy.datasets.pbmc3k()``.

Samples: 2,700 Features: 32,738 Sparsity: 97.4% Missing values: No

Preprocessing: `normalize_total`

Raw UMI counts (`profile.json`: mean ~2367 total counts per cell) have very different per-cell totals; `DATA_DESCRIPTION.md` notes no normalization has been applied yet. Library-size normalization removes sequencing-depth differences between cells before comparing expression profiles.

Preprocessing: `log1p`

Raw counts are heavily right-skewed; `DATA_DESCRIPTION.md` notes a small number of highly-expressed genes can dominate raw totals. Log transform compresses this skew so distance-based methods (PCA, UMAP) aren't dominated by a handful of extreme values.

Preprocessing: `select_hvg`

32,738 genes is far more than 2,700 cells, and 97.4% of entries are zero (`profile.json`), so most genes carry little signal. Selecting the top 2,000 highly variable genes keeps the informative signal and keeps distance computations from being dominated by noise genes.

Methods and hyperparameters

pca (role: general_purpose, status: reused)

Hyperparameters: `n_components=50`

Required linear baseline. 50 components is the standard scRNA-seq choice: enough to capture most biological variance across 2,000 genes without overfitting to per-gene noise given only 2,700 cells.

umap (role: visualization_only, status: reused)

Hyperparameters: `n_neighbors=15`, `min_dist=0.1`, `n_components=2`, `random_state=0`


Standard nonlinear manifold method for scRNA-seq visualization. `n_neighbors=15` and `min_dist=0.1` are UMAP's own defaults; with 2,700 cells there's no strong reason yet to move off the defaults, this is a first-pass plan for pipeline testing, not a final analysis.

Quantitative metrics

method	trustworthiness	sum(EVR)	var. used	PCs for 90%	MDS stress	label silhouette
pca	0.9344	0.2904	0.2904	200		
umap	0.7656					


n=2,700, showing dims 1-2 of 50

pbmc: pca



n=2,700

pbmc: umap



Limitations

No ground-truth labels available for this dataset; embedding quality assessed only via label-free trustworthiness, no supervised sanity check possible.

No adversarial critique/revision pass was performed for this run (critic subagent not yet built).